

# MLOps Project Report

Yunhang Chi, Wenxi Huang, Shiyu Mao, Zhongke Sun

April 10, 2025

## 1 Introduction

This report presents an MLOps pipeline built around the classic iris classification problem. The goal of this project was to demonstrate a practical approach to model training, evaluation, deployment, and basic versioning within a simplified Python ecosystem. Although the iris dataset is relatively small and straightforward, it offers a convenient framework for exploring key concepts like reproducible training, containerization, and command-line automation.

## 2 Workflow Design

A typical workflow begins by calling the training step. The user can install dependencies from `requirements.txt` and then invoke `python model.training.py` to produce a trained random forest model. If they wish to evaluate performance again or modify certain parameters, they can adjust the training script and re-run the process. For inference, `python model.prediction.py` loads the saved model and displays predictions for a sample input. Once satisfied with local results, the user can containerize the project. The project includes a `Dockerfile` that sets up a minimal Python environment, installs the required libraries, and defaults to running the model code upon container startup. With the Docker image built, `CLI.py` offers straightforward commands like `build`, `push`, `pull`, and `run_model` for sharing or launching the container.

### 3 Project Structure

The project follows a modular design to keep training, evaluation, inference, and data preprocessing tasks both transparent and extensible. Each stage is defined in a separate Python script:

**preprocessing.py** Loads and prepares the iris dataset for training and testing. It also includes an optional exploratory analysis function that generates basic visualizations such as scatterplots and correlation matrices.

**model\_training.py** Implements a `train_model()` function that trains a random forest classifier, prints training duration, and saves the final model to a local file (`model.pkl`). Once the model is trained, the script triggers an evaluation routine to measure classification performance.

**model\_evaluation.py** Defines `evaluate_model(...)`, which computes metrics like accuracy, precision, recall, and F1 score. It also plots and saves a confusion matrix. The script includes per-class accuracy reports, serving as a basic illustration of how fairness or bias checks might be performed in a more complex scenario.

**model\_prediction.py** Loads the persisted model file and demonstrates a simple inference process. In this example, it predicts the species of a single test instance, though the logic can be adapted to support batch or interactive prediction.

**model\_registry.py** Illustrate a simple versioning workflow. Each time the `register_model()` function is called, it increments the version by counting existing “Version” lines in the registry log, logs metadata (path and timestamp) to `model_registry.txt`, and appends a “Version X” header and a separator line for clarity.

**CLI.py** Handles both container operations and local model tasks. Users can build, run, push, or pull a Docker image, and can also train or predict locally without invoking Docker.

**get\_path.py** Defines relevant functions to dynamically generate and manage project directories, such as saving models, logs, and plots outputs.

### A Appendix: Screenshots

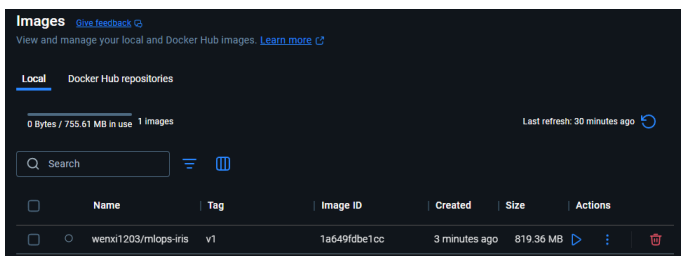


Figure 1: Build




Figure 2: Run Container




Figure 3: Push Docker Hub




Figure 4: Delete Figure and Pull

```

C:\Users\hwx1\Desktop\MLops\docker_display\docker_deploy 2\docker_deploy>python CLI.py run_model
Model training completed in 0.07 seconds.
Evaluation Metrics:
Accuracy : 1.00
Precision: 1.00
Recall : 1.00
F1 Score : 1.00
Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
Confusion matrix saved as confusion_matrix.png
Per-Class Accuracy:
Class 0 accuracy: 1.00
Class 1 accuracy: 1.00
Class 2 accuracy: 1.00
Model saved as model.pkl
Registered model version 1
Log saved to: C:\Users\hwx1\Desktop\MLops\docker_display\docker_deploy 2\docker_deploy\model_registry.txt
Model loaded from model.pkl
Prediction for [[6.1, 2.8, 4.7, 1.2]]: [1]
true label: 1

```

```

C:\Users\hwx1\Desktop\MLops\docker_display\docker_deploy 2\docker_deploy>python CLI.py train
Model training completed in 0.28 seconds.
Evaluation Metrics:
Accuracy : 1.00
Precision: 1.00
Recall : 1.00
F1 Score : 1.00
Confusion Matrix:
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
Confusion matrix saved as confusion_matrix.png
Per-Class Accuracy:
Class 0 accuracy: 1.00
Class 1 accuracy: 1.00
Class 2 accuracy: 1.00
Model saved as model.pkl
Registered model version 1
Log saved to: C:\Users\hwx1\Desktop\MLops\docker_display\docker_deploy 2\docker_deploy\model_registry.txt

```

Figure 5: Run and Train the Model

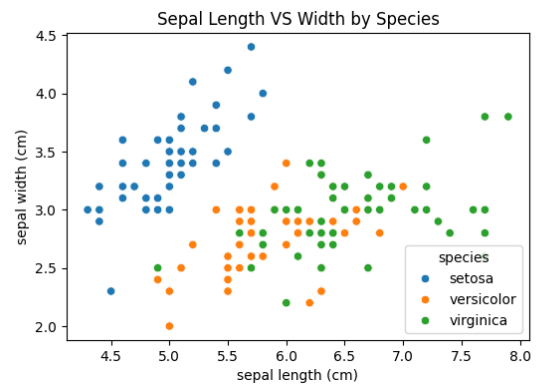
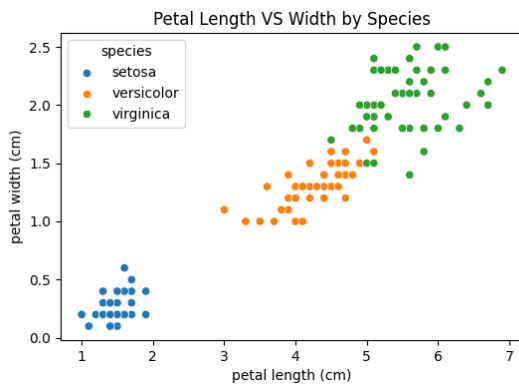
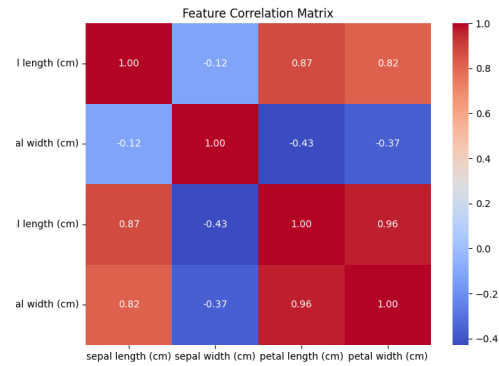
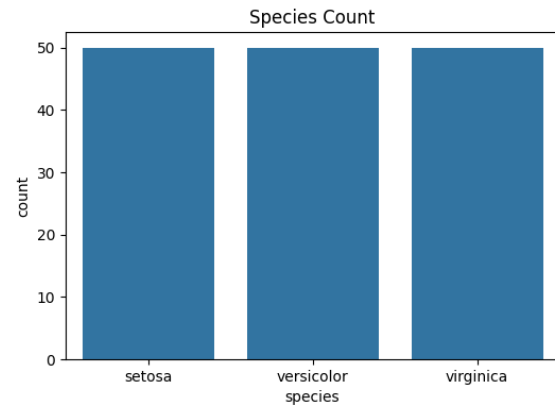


Figure 6: Visualization